

Research Roadmap

Institutional Independence — from concept to standard

This is a research program, not a finished paper. The conceptual and operational layers are built; the validation layers require deployment data and institutional partners. The roadmap below marks what is done, what is in progress, and where collaboration is the gating constraint.




Figure 1: Six phases from concept to standard. Phases 1–3 (conceptual framework, operational metrics, reference profiles) are built and reproducible. Phase 4 (deployment-log validation) is next and is gated on access to a real recommendation log. Phases 5–6 (cross-organizational benchmark, standardization) need deployment data and institutional partners.

The arc

The program moves through six phases, each one making institutional independence more *real* — from a defined property, to a measurable one, to an auditable one, to a standardized one.

Phase	What it delivers	Status	What it proves / what it needs
1 • Conceptual framework	Institutional independence defined as a governance property orthogonal to robustness, fairness, transparency; the confirmation-loop mechanism	Done	<i>Proves:</i> the property is distinct and not a relabeling. → position paper
2 • Operational metrics	Five log-computable measures (incentive sensitivity, diversity, dissent robustness, counterfactual consistency, escalation dependence) + reference implementation	Done	<i>Proves:</i> the property is operationalizable, not just conceptual. → <code>metrics_skeleton.py</code>

Phase	What it delivers	Status	What it proves / what it needs
3 · Reference profiles	The metrics applied to three public governance structures (OpenAI SAG, Anthropic CAI, Meta OB), each score citation-backed; standards-gap analysis (NIST / ISO 42001 / EU AI Act)	Done	<i>Proves:</i> the framework reads real structures and fills a documented standards gap. → reference profiles, standards gap
4 · Deployment log validation	Run the five measures on a real advisor's recommendation log over time; establish whether institutional drift occurs in practice, and at what selection strength	Next	<i>Needs:</i> access to a deployment log (internal governance function, or research partnership). This is the gating step — and the one collaboration unlocks.
5 · Cross-organizational benchmark	The same fixed cases run across institutions with different incentive structures, isolating the institution's contribution from the model's; a shared Independence Benchmark	Future	<i>Needs:</i> multiple participating organizations; a neutral host (a lab, a standards body, an auditor).
6 · Standardization proposal	Institutional independence proposed as an evaluable property for inclusion in AI governance standards (a NIST profile, an ISO 42001 control, an EU AI Act conformity criterion)	Future	<i>Needs:</i> phases 4–5 evidence; engagement with a standards body. The endpoint: independence becomes an auditable requirement, not a paper.

Where it stands

Built (Phases 1-3): the framework is conceptually complete, operationally specified in code, and validated against real public structures. Everything needed to *define and measure* institutional independence exists and is reproducible. A reviewer can read the property, run the metrics, and see the profiles.

The frontier (Phase 4): the open question is *prevalence* — does institutional drift actually occur in deployed advisory systems, and how strongly? This cannot be answered from public documentation or simulation alone. It requires the one thing the program does not yet have: **a real deployment log to run the protocol against.** The simulation (`sim_loop.py`) shows the mechanism is *sufficient* to produce drift; Phase 4 tests whether it is *present*.

The horizon (Phases 5-6): if drift is real and measurable across organizations, the natural endpoint is standardization — institutional independence named as an evaluable property in the

frameworks that currently lack it (the standards-gap analysis identifies exactly where it would slot in).

The collaboration hook

The program is deliberately structured so that its gating step is also its most valuable contribution. **Phase 4 needs a deployment log; whoever provides one gets the first real Independence Profile of their own advisory system.** That is the natural point of partnership:

- **A frontier lab** with a safety-review or release-advisory pipeline could run the protocol on its own logs — the first empirical test of whether its oversight stays independent under pressure.
- **A regulator or auditor** could pilot the Independence Benchmark as a deployment-log audit method, ahead of the standards that will eventually require one.
- **An organizational-theory or AI-governance researcher** could co-author the Phase 4–5 empirical work, vetting the framework’s grounding in the organizational literature it draws on.

The framework is built to be handed off, not held. Its value is realized when someone runs it on a system that matters.

*Framework, code, profiles, and standards analysis: github.com/ssloves/institutional-independence
Se Mi Song — ssloves.github.io · songsemisong@gmail.com · open to collaboration on Phase 4 deployment-log validation.*