

Institutional Independence in Practice

Applying the Framework to Three Public AI Governance Structures

Se Mi Song · companion case study to *The Institutional Confirmation Loop*

This note applies the institutional-independence lens to three publicly documented AI governance structures. It is a demonstration of the framework as an analytical tool, not an audit or a claim about any organization’s conduct. Each analysis uses only public documentation and asks one question: where in this structure would the framework direct an auditor to look?

Why these three

Institutional independence is the capacity of an oversight function to track objective risk despite the incentives of the organization it sits inside. The three structures below were chosen because each instantiates the advisor-institution relationship differently — together they span a useful range, from advisory-only to binding-but-bounded.

Structure	The advisor	Independence design
OpenAI — Safety Advisory Group	internal expert body advising on release	advisory; leadership decides
Anthropic — Constitutional AI	the model critiquing its own outputs	principle-based self-revision
Meta — Oversight Board	external body reviewing content decisions	binding on cases, advisory on policy

Case 1 — OpenAI’s Safety Advisory Group (SAG)

The structure. Under OpenAI’s publicly published Preparedness Framework, a cross-functional Safety Advisory Group reviews capability and safeguard reports and makes deployment recommendations; OpenAI Leadership makes the final call, with board-level oversight. The framework states plainly that leadership can act without the SAG’s participation — the SAG “does not have the ability to filibuster.”

What the framework sees. This is a textbook advisor-institution configuration, and the public documentation already surfaces the two variables the framework cares about. First, **Escalation Dependence and the decision right**: an advisor that recommends but cannot block has, by design, an independence ceiling — its influence is mediated entirely by whether the institution chooses to act on it. The framework would direct an auditor to the *gap between recommendation and decision*: how often, and under what conditions, do recommendations not survive to the deployment decision? Second, and more directly, the framework’s published competitive-adjustment provision — that safeguards may be revisited if a competitor releases comparable capability — is an explicit **Institutional Incentive Sensitivity** surface. It writes a specific institutional incentive (competitive position) into the safety assessment itself. The framework does not judge

whether this is right; it identifies it as exactly the variable *Counterfactual Consistency* would hold the evidence against: same capability and risk evidence, different competitive context — does the safeguard threshold move?

Where to look. The Independence Profile here is dominated by the recommendation→decision gap (Escalation Dependence) and the competitive-adjustment clause (Incentive Sensitivity). Both are measurable from release records over time.

Case 2 — Anthropic’s Constitutional AI

The structure. Constitutional AI trains a model to critique and revise its own outputs against a written set of principles (the “constitution”), reducing reliance on human labels. The model itself is, in effect, an advisor on its own behavior — and the principles are author-chosen.

What the framework sees. This is the most interesting case, because Constitutional AI is partly an *attempt to build independence in*. The self-critique step gives the model a standard external to any single output to judge against — structurally similar to what institutional independence asks for: an assessment anchored to stated principles rather than to whatever is locally convenient. In that sense CAI internalizes a form of the property. But the framework also locates the residual question precisely: **the constitution is chosen by the institution**. Independence from the *institution’s incentives* is only as strong as the principles’ independence from those incentives — so the framework would direct attention not to the self-critique mechanism (which is sound) but to the *governance of the constitution itself*: who sets and revises the principles, and whether that process is insulated from product or commercial pressure. Here the framework’s value is to separate two things often conflated — independence of *the model’s reasoning* (CAI provides this) from independence of *the standard the reasoning serves* (a governance question CAI does not by itself settle).

Where to look. The Independence Profile shifts up a level: not the model’s per-output consistency (which CAI strengthens) but the *constitutional amendment process* as the locus where institutional incentive could enter.

Case 3 — Meta’s Oversight Board

The structure. An external board reviews content-moderation decisions. Its case decisions are binding on Meta; its broader policy recommendations are not. Independent reporting notes the board has overturned a large majority of the decisions it reviewed — but also that Meta retains control over the scope and precedent of the board’s work.

What the framework sees. This case shows independence is not binary but **domain-specific** — which is exactly why the framework reports a *profile*, not a single score. On binding case decisions, the board scores high on the property: the high overturn rate is direct evidence of *Dissent Robustness* — an advisor that frequently rules against the institution that created it is demonstrably surfacing inconvenient conclusions. But on policy recommendations (non-binding) and on jurisdiction (institution-controlled), the same body has a lower independence ceiling. The framework makes this legible: a single structure can be strongly independent on one axis (binding cases) and weakly independent on another (policy influence, scope), and an honest Independence Profile records both rather than averaging them into a misleading single number.

Where to look. Two very different cells of the profile: high *Dissent Robustness* on binding cases, lower independence on the non-binding/jurisdictional axes where the institution retains control.

What the three cases show together

	Strongest independence	Framework directs attention to
OpenAI SAG	expert review quality	recommendation→decision gap; competitive-adjustment clause
Anthropic CAI	per-output reasoning	governance of the constitution itself
Meta OB	binding case decisions	non-binding policy + institution-controlled scope

Three takeaways for anyone designing or evaluating AI oversight:

1. **Independence has a ceiling set by the decision right.** An advisor that recommends but cannot bind (SAG) is independence-limited by construction; one that binds on some matters but not others (OB) is independent in patches. The decision architecture, not the reviewers’ competence, sets the ceiling.
2. **Independence of reasoning $\neq$ independence of the standard.** Constitutional AI shows you can make an advisor reason against a fixed standard *and still* leave open who controls the standard. Both must be governed.
3. **Independence is a profile, not a verdict.** Every one of these structures is strong on one axis and bounded on another. The framework’s job is not to grade them but to make the axes — and the places institutional incentive can enter — explicit and auditable.

This is what the framework is *for*: not to declare an oversight body captured or independent, but to turn “is this oversight real?” into a small set of specific, checkable questions about where a given structure’s independence is strong, where it is bounded, and where an institution’s incentives have a documented path into the assessment.

A reference evaluation: filling in the profile

To show what the framework *produces*, we render each structure as an Independence Profile across the five measures. **An essential caveat: these are structural estimates inferred from public documentation, not measurements.** A real Independence Profile requires deployment logs and the perturbation protocol (fix evidence, vary incentive); we have neither here. What follows shows the *shape of the framework’s output* — how a completed profile reads — using each structure’s published design to esti-

mate where its independence is strong or bounded. The honest version of this exercise on real logs is the empirical program the paper calls for.

Independence Profiles — structural assessment from public documentation (illustrative, not measured)




Figure 1: Independence Profiles for the three structures, estimated from public documentation. Filled circles indicate stronger independence on that axis (less structural room for institutional incentive to enter). These are illustrative structural assessments, not measurements from deployment logs.

How to read the estimates (each is a structural inference with a one-line public-document basis):

- **OpenAI SAG — Escalation Dependence (low):** the advisor recommends but cannot bind; published framework states the SAG cannot “filibuster.” The decision right sits entirely downstream, which structurally caps independence on this axis. *Incentive Sensitivity* is also bounded by the explicit competitive-adjustment provision, which routes an institutional incentive into the safeguard threshold by design.
- **Anthropic CAI — Counterfactual Consistency (high):** self-critique against a fixed written standard is structurally well-suited to holding an assessment steady across framings. The estimate is *bounded above*, not maxed, because the standard itself (the constitution) is institution-set — the residual question the profile flags rather than resolves.
- **Meta OB — Dissent Robustness (high):** the board has publicly overturned a large majority of reviewed decisions, direct evidence of an advisor surfacing conclusions against its sponsoring institution. *Escalation Dependence* is lower because policy recommendations are non-binding and jurisdiction is institution-controlled.

The point is not the specific circle counts — reasonable analysts would place some differently — but that the framework forces each judgment to attach to a *named axis* and a *documentable basis*, and records strength and limitation side by side rather than collapsing them into a verdict. That is what makes an Independence Profile an audit artifact rather than an opinion.

Framework, metrics, and simulation: github.com/ssloves/institutional-independence · All three structures analyzed from public documentation (OpenAI Preparedness Frame-

work v2; Anthropic, “Constitutional AI” and “Claude’s Constitution”; Meta Oversight Board charter and transparency reports).