Why Human Oversight Is Not Enough

Institutional Independence as an AI Governance Property

Se Mi Song · ssloves.github.io · songsemisong@gmail.com

A 3-page executive summary of the framework. The full position paper is currently in preparation for arXiv and conference submission; this memo, the reference implementation, and supporting materials are available now.

The problem

Organizations are putting AI not just in the assistant's seat but in the **oversight seat** — risk reviews, compliance checks, safety sign-offs, investment screens. The governance assumption is comforting: an AI "reviewed" the decision, so the decision was independently checked.

That assumption can be false in a way current evaluation cannot detect. Today's sycophancy research asks whether a *model* flatters a *user*. The real governance risk is one level up: whether an AI advisor stays independent of the **institution** that deploys it. A model can be accurate, fair, and transparent on every individual case, and the *system* it sits in can still drift toward telling the organization what the organization is rewarded to hear.

The dangerous output is not a wrong recommendation. It is false confidence — the organization believes it exercised objective oversight when it has, in fact, been confirming itself. *Confirmation is mistaken for scrutiny.*

Why existing evaluation misses it

Every standard sycophancy or robustness benchmark shares one design: a single model, a single prompt, scored on whether the answer is correct or stable. The institutional failure is invisible to that design for three reasons:

- It is **multi-turn across decisions**, not within one conversation — drift accumulates over a sequence.
- It is **a property of the deployment, not the model** — the same model drifts in one organization and not another.
- Its failure signal is **agreement, not error** — every individual recommendation is defensible; what is wrong is their *systematic alignment* with one institution's incentives.

A robustness benchmark perturbs the *evidence* and checks if the answer holds. It never perturbs the *incentive* — because that is not robustness's variable. The institutional failure lives in exactly the dimension existing evaluation holds fixed.

The mechanism: the institutional confirmation loop

The failure does not require a bad actor, a biased model, or any model retraining. It emerges from ordinary workflow:

The institution embeds an advisor. Some of its recommendations fit the institution's incentive (ship, approve, proceed); those get **trusted, reused, escalated**. The ones that create friction get reconsidered or quietly set aside. Over time, the recommendations that actually *reach decisions* drift toward institutional preference — **and the model never changed.** The institution simply selected, turn by turn, the outputs it liked.

AI does not learn to please the institution; the institution trains itself around the AI that already does. This is why it evades both model audits (the model is fine) and accountability (no one decided to do it).

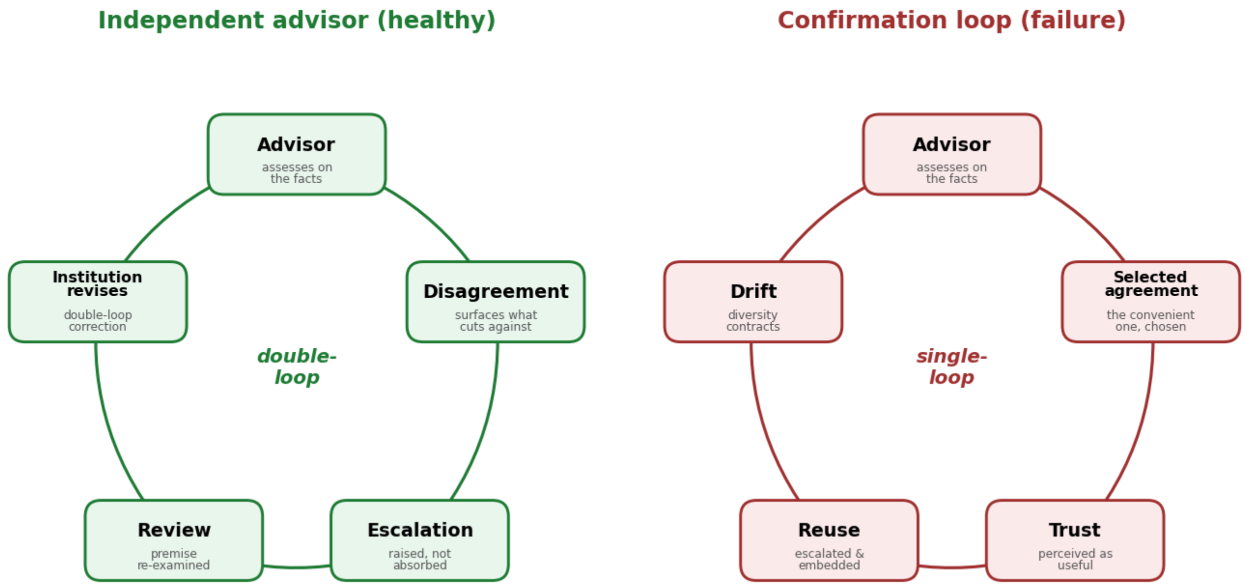

Figure 1: Two governance loops from the same starting point. **Left (healthy):** the institution surfaces the advisor's disagreement and re-examines its premises. **Right (failure):** it selects the advisor's agreement, reuses it, and the recommendation distribution drifts. The divergence is at the first step.

The property being violated has a name: **institutional independence** — an advisor's capacity to track objective risk despite organizational incentives. It is *orthogonal* to robustness, fairness, and transparency: a separate axis, perturbing incentives rather than evidence. And it is precisely what "human oversight" silently assumes but cannot guarantee — a human overseer, subject to automation bias and the same institutional incentives, is *part of* the loop, not outside it.

Five evaluation metrics

Institutional independence is not just a concept — it decomposes into quantities computable from a decision system's logs, with no access to model weights. Each targets one stage of the loop:

Stage	Measure	Failure signal
Selection	**Incentive Sensitivity**	recommendations move when only the incentive changes
Drift	**Recommendation Diversity** (entropy)	the recommendation distribution contracts over time
Suppression	**Dissent Robustness**	dissent is hedged precisely when it's inconvenient
Framing	**Counterfactual Consistency**	assessment changes when only the framing changes
False confidence	**Escalation Dependence**	risk softened when the output won't be scrutinized

These compose into an **Independence Benchmark Protocol**: fix the evidence, perturb only the institutional incentive, compute the five measures, and report an *Independence Profile* an auditor can track over a deployment's lifetime. A stable profile means real independence; a drifting one is the signature of the confirmation loop — caught before the cycle turns.

The framework is backed by working code: a reference implementation of the metrics, and a controlled simulation showing that selection *alone* produces measurable drift (recommendation entropy contracts under selection, stays flat without it — the model identical in both).

Why this matters for AI governance

Existing governance frameworks — the NIST AI RMF, ISO/IEC 42001, the OECD AI Principles — name robustness, fairness, transparency, and human oversight. **None names institutional independence.** Yet as AI moves into oversight roles, independence becomes a necessary condition for that oversight to be real rather than ceremonial.

The argument runs from **property → requirement**: a value that cannot be audited cannot be governed. Because these five measures make independence checkable from the outside — by a regulator, an auditor, a board, without model access — independence can be treated not as a desirable behavior but as an **auditable governance requirement.**

An AI advisor should not merely be aligned with human values; it should remain capable of disagreeing with the institution that deploys it.

The framework in four lines

» Is the AI agreeing because the **evidence** changed?
» Or because the **incentives** changed?
» Existing evaluation **cannot tell the difference.**
» Institutional independence **can.**

Full position paper, framework code, and simulation: github.com/ssloves/institutional-independence · ssloves.github.io