

Reference Independence Profiles

Three public AI governance structures, scored with cited evidence

*Companion to the Institutional Independence framework. Each score is a **structural assessment from public documentation**, not a measurement from deployment logs. Real profiles require the perturbation protocol (fix evidence, vary incentive) run on actual logs. What this document adds: every score carries an explicit, citable basis, so the profile is auditable rather than asserted.*

Scale: ● = structural support for independence on this axis; ○ = structural room for institutional incentive to enter. Five circles per axis.

OpenAI — Safety Advisory Group (SAG)

Measure	Profile	Basis (public documentation)
Incentive Sensitivity	●●○○○	Preparedness Framework v2 includes an explicit competitive-adjustment provision: safeguards may be revisited if a competitor releases comparable High/Critical capability. An institutional incentive (competitive position) is written into the safeguard threshold by design. [1]
Recommendation Diversity	●●●○○	SAG reviews Capabilities and Safeguards Reports and can recommend a <i>range</i> of outcomes (approve / request further evaluation / require stronger protections), preserving output diversity at the recommendation stage. [1]
Dissent Robustness	●●○○○	The SAG can raise objections, but the framework states leadership may proceed without SAG participation — “the SAG does not have the ability to filibuster.” Dissent is structurally non-binding. [1][2]
Counterfactual Consistency	●●○○○	No public mechanism fixes the assessment against institutional framing; the competitive-adjustment clause explicitly permits the threshold to move with competitive context. [1]

Measure	Profile	Basis (public documentation)
Escalation Dependence	●○○○○	The decision right sits entirely with leadership; the advisor recommends but cannot bind. Independence on this axis is capped by construction. [1][2]

Reads as: strong technical review, structurally bounded authority. The framework directs an auditor to the recommendation→decision gap and the competitive-adjustment clause.

Anthropic — Constitutional AI (CAI)

Measure	Profile	Basis (public documentation)
Incentive Sensitivity	●●●○○	Self-critique is anchored to written principles rather than case-by-case institutional preference, which structurally dampens incentive sensitivity at the output level. [3][4]
Recommendation Diversity	●●●○○	The model revises against principles drawn at random across many principles; the process is not collapsed to a single objective. [3]
Dissent Robustness	●●●○○	CAI is explicitly designed to engage rather than evade — to <i>explain objections</i> to harmful requests rather than refuse flatly — preserving a form of dissent. [4]
Counterfactual Consistency	●●●●○	Critique against a <i>fixed written standard</i> is structurally well-suited to holding an assessment steady across framings. Bounded above (not ●●●●●) because the standard itself is institution-set. [3][4]
Escalation Dependence	●●●○○	The constitution is applied uniformly in training rather than conditioned on a downstream audience; but the property lives one level up, in who governs the constitution. [4]

Reads as: independence of *reasoning* partly built in; the open question is independence of the *standard*. The framework directs attention to the constitution’s amendment governance, not the self-critique mechanism.

Meta — Oversight Board (OB)

Measure	Profile	Basis (public documentation)
Incentive Sensitivity	●●●○○	The board grounds decisions in international human-rights standards external to Meta's commercial incentives, dampening institutional-incentive sensitivity on case decisions. [5]
Recommendation Diversity	●●●●○	The board issues a wide range of outcomes — overturn, uphold, mixed, plus policy recommendations — across a large and growing caseload. [5][6]
Dissent Robustness	●●●●●	Independent analyses report the board has overturned roughly 80% of more than 100 reviewed decisions — direct, quantified evidence of an advisor ruling against its sponsoring institution. [6]
Counterfactual Consistency	●●●○○	Human-rights grounding supports consistent reasoning across cases, but the board's <i>jurisdiction</i> is set by Meta, bounding what it can hold constant. [5][7]
Escalation Dependence	●●○○○	Case decisions are binding ; policy recommendations are not (Meta responds within 60 days; ~74% of 300+ recommendations implemented or progressed). Independence is high on binding cases, lower on policy and scope. [7][8]

Reads as: the strongest binding-case independence of the three (high Dissent Robustness), bounded on the non-binding/jurisdictional axes. A clear illustration of why independence is a *profile*, not a verdict.

What the three profiles show

	Strongest axis	Bounded axis	The lesson
OpenAI SAG	recommendation range	escalation / decision right	independence capped by the decision architecture

	Strongest axis	Bounded axis	The lesson
Anthropic CAI	counterfactual consistency	governance of the standard	reasoning-independence $\neq$ standard-independence
Meta OB	dissent robustness (binding cases)	escalation (policy non-binding)	independence is domain-specific, not global

No structure is uniformly strong; each is independent on some axes and bounded on others. **That is the framework’s contribution in one view: it replaces “is this oversight real?” with a per-axis, evidence-cited profile that records strength and limitation side by side — and shows exactly where an institution’s incentives have a documented path into the assessment.**

Sources

[1] OpenAI, *Preparedness Framework v2* (April 2025) — Safety Advisory Group structure, decision rights, competitive-adjustment provision. [2] OpenAI, *Safety practices / safety update* — SAG reviews, leadership final decisions, board oversight. [3] Bai et al., *Constitutional AI: Harmlessness from AI Feedback*, arXiv:2212.08073 (2022) — two-phase self-critique/revision process. [4] Anthropic, *Claude’s Constitution* — principles, application in training, sources of principles. [5] Meta Oversight Board — charter and mission; human-rights grounding; case-review scope. [6] Ang & Haristya (2024), *The Governance, Legitimacy and Efficacy of Facebook’s Oversight Board* — ~80% overturn rate across 100+ decisions. [7] Meta Transparency Center — binding case decisions vs. non-binding recommendations; 60-day response commitment; jurisdiction set by Meta. [8] Meta Oversight Board, *2024 Annual Report* — 300+ recommendations since 2021; ~74% implemented or progressed.

All figures are drawn from public documentation as of mid-2026 and may change. Scores are structural estimates for illustrating the framework’s output, not measurements of organizational conduct.