

The Standards Gap

Why No Existing AI Governance Standard Evaluates Institutional Independence

Se Mi Song · companion analysis to *The Institutional Confirmation Loop*

This note asks one question of three leading AI governance standards: can they detect whether an AI advisor remains independent of the institution deploying it? The answer, in each case, is no — and the reason is the same. Each standard constrains the system; none evaluates the system-institution relationship. This is the gap institutional independence fills. All claims below cite public standard text.

The one question

Institutional independence is the capacity of an AI advisor to track objective risk despite the incentives of the organization it sits inside. Evaluating it requires one specific move: **perturb the institutional incentive while holding the evidence fixed, and measure whether the assessment moves** (the perturbation that defines the property). The question for each standard is simply: does it require, or even enable, that move?

NIST AI RMF (AI 100-1)

What it covers. Four functions — Govern, Map, Measure, Manage — aligned to seven trustworthy-AI characteristics: valid and reliable, safe, secure and resilient, accountable and transparent, explainable and interpretable, privacy-enhanced, and fair with harmful bias managed. [1][2]

Can it evaluate institutional independence? No.

- **None of the seven characteristics is the advisor-institution relationship.** Robustness is resistance to input perturbation; fairness is equitable treatment across subject groups; transparency is legibility of reasoning. Independence — resistance to *incentive* perturbation — is not among them, and does not reduce to any of them.
- **Govern presumes the independent overseer it cannot verify.** The practical sign that Govern is working, in NIST’s own framing, is that “one person or committee can stop or approve an AI use case” using documented criteria. [3] But that is exactly the human overseer the confirmation loop captures: the standard assumes the approver is independent; it provides no measure of whether they are. Govern institutionalizes oversight; it does not test the oversight’s independence from institutional incentive.
- **Measure perturbs evidence, not incentives.** The Measure function develops quantitative methods to assess risks and track drift — but “drift” here means data/performance drift, not *governance* drift (contraction of an advisor’s recommendation distribution toward institutional preference). No subcategory varies the institutional principal with evidence held fixed.

Verdict: NIST AI RMF names what to govern and tells you to have an overseer. It does not, and structurally cannot, tell you whether that overseer has stayed independent.

ISO/IEC 42001:2023

What it covers. A certifiable AI Management System (AIMS) — the management-system controls (policies, roles, processes, documentation, continual improvement) that wrap an organization’s AI, in the ISO 27001 / management-system tradition. [4]

Can it evaluate institutional independence? No.

- **It governs the management process, not the assessment's independence.** ISO 42001 certifies that an organization *has* a system for managing AI risk — roles assigned, processes documented, improvement tracked. It is procedural by design. An organization can be fully 42001-certified while its AI advisor's recommendations drift toward institutional preference, because nothing in the management-system checklist measures the *content* of the advisor-institution relationship over time.
- **Certification is a conformity check, not an incentive-perturbation test.** The audit confirms the system exists and operates as documented. It does not present the advisor with fixed evidence under varied institutional incentive and check whether the assessment holds — the one test independence requires.

Verdict: ISO 42001 verifies you have a governance process. Institutional independence is a property of what that process *produces* under incentive pressure — which a management-system audit does not reach.

EU AI Act (Regulation 2024/1689)

What it covers. Binding, risk-tiered law. For high-risk systems: risk management (Art. 9), data governance (Art. 10), technical documentation (Art. 11), record-keeping (Art. 12), transparency (Art. 13), **human oversight (Art. 14)**, and accuracy/robustness/cybersecurity (Art. 15), verified through conformity assessment. [5][6]

Can it evaluate institutional independence? No — and it is the most instructive case, because it comes closest and still misses.

- **Article 15 is evidence-perturbation, explicitly.** It requires resistance to “errors, faults, and manipulation” and adversarial robustness. [6] That is the *evidence* axis — perturb the input, does the output hold. It says nothing about perturbing the *institutional incentive*.
- **Article 14 names automation bias — then hands it to the individual.** Remarkably, the Act explicitly requires that human overseers “remain aware of the possible tendency of automatically relying or over-relying on the output” — automation bias — “in particular for systems used to provide recommendations for decisions to be taken by natural persons.” [5] This is precisely the loop's mechanism. But the Act's remedy is an *individual* competence (“remain aware,” retain the ability to override). It does not address the *institutional* version: an overseer who is individually aware but organizationally incentivized to defer, embedded in workflows that select for the convenient recommendation. The Act regulates the operator; the confirmation loop operates on the institution around the operator.
- **Conformity assessment is one-time and structural, not longitudinal.** Independence fails *over time*, as selection accumulates. A pre-market conformity assessment (even a third-party one) checks the system at a point; it does not track the recommendation distribution across a deployment's life. The Act's record-keeping (Art. 12) *generates the logs* a longitudinal independence audit would need — but the Act does not specify that audit.

Verdict: The EU AI Act regulates human oversight and even names automation bias, but locates both at the individual operator. The institutional confirmation loop is the gap *between* the aware individual and the incentive-shaped organization — which the Act's structure does not reach.

The common shape of the gap

Standard	What it constrains	Why it misses independence
NIST AI RMF	the system’s trustworthy characteristics	none of the seven is the advisor-institution relationship; Govern presumes an independent overseer it cannot verify
ISO/IEC 42001	the management process around AI	certifies a process exists; does not test what the process produces under incentive pressure
EU AI Act	high-risk system requirements + human oversight	regulates the individual operator; names automation bias but not its institutional aggregation

Every one constrains the system so the institution can trust it. None constrains the relationship between the system and the institution — which is the one thing that fails when oversight is captured. That inversion is exactly what institutional independence supplies, and why it belongs in the standards rather than beneath them.

This is the framework’s claim to standards relevance: it is not a competitor to NIST, ISO, or the EU AI Act. It is the **missing axis** each of them would need to detect a confirmation loop — a per-axis, log-computable measure of whether an AI advisor, placed inside an institution’s oversight, has stayed independent of that institution’s incentives. The standards generate the logs (EU AI Act Art. 12 record-keeping; NIST Measure) and mandate the overseer (NIST Govern; EU AI Act Art. 14). What they lack is the test. This framework is that test.

Sources

[1] NIST, *AI Risk Management Framework (AI RMF 1.0)*, NIST AI 100-1 (Jan 2023) — four functions; seven trustworthy-AI characteristics. [2] NIST AIRC — AI RMF Core (Govern, Map, Measure, Manage). [3] NIST AI RMF Govern function — oversight as “one person or committee can stop or approve an AI use case” with documented criteria. [4] ISO/IEC 42001:2023 — AI Management System (AIMS); certifiable management-system standard. [5] EU AI Act (Reg. 2024/1689), Article 14 — human oversight; explicit automation-bias provision for recommendation systems. [6] EU AI Act, Article 15 — accuracy, robustness against errors and manipulation, adversarial resistance; Article 12 — record-keeping/logs.

Standard texts as of mid-2026; implementation timelines and revisions (e.g., NIST AI RMF 1.1, EU AI Act high-risk enforcement) are evolving. Claims describe each standard’s scope, not its adequacy in other respects.